An Intro to Machine Learning

Rouhollah Abolhasani

What is Machine Learning?

Speech Recognition

- Alexa and Siri
  - Extract vowels and consonants from the signals
  - Recognize Words, Sentences, and Sentiment
  - Act accordingly.

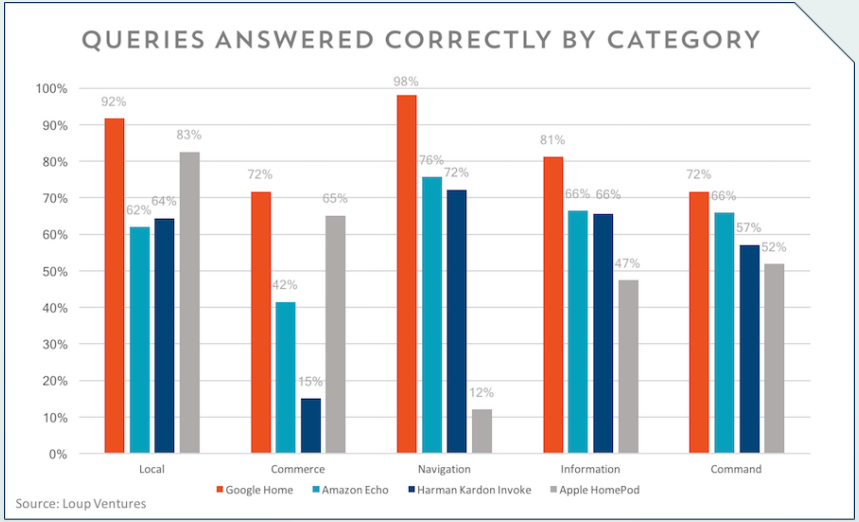

Time-series Forecasting

- Time-series
  - Stock Market
  - Texts, Tweets, ...
- COTA: Customer Obsession Ticket Assistant
  - Handling Uber support tickets
  - Processing tickets and proposing solutions

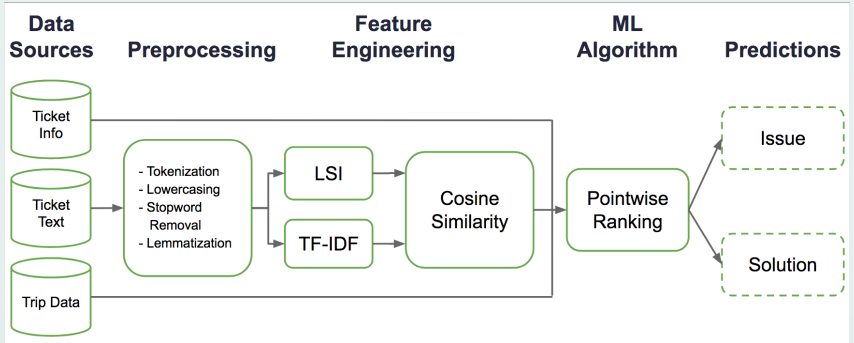

https://eng.uber.com/cota/

Classification

- Image Classification
  - Cat vs. dog
  - Image segmentation
- Old-fashioned Spam Detector
  - What would we do?
  - Writing rules
  - More rules
  - And more rules …
- Better Spam Detector
  - Automatic feature extraction
  - Model training
  - Classifying new emails

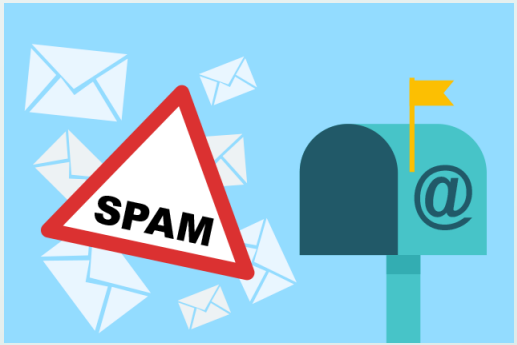

Rule-based vs. Learning

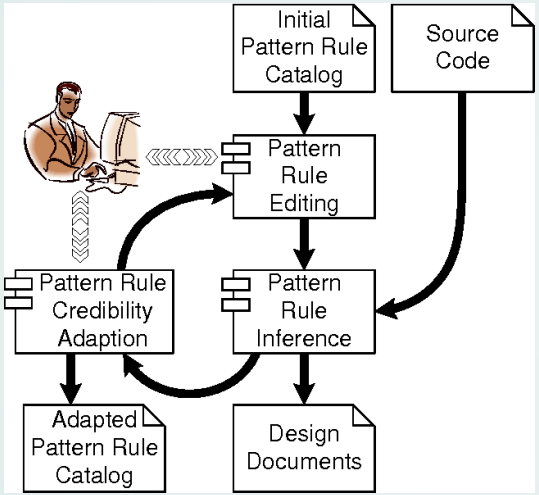

Rule –based Programming

Write some rules. Evaluate them. Add some again, and again, …

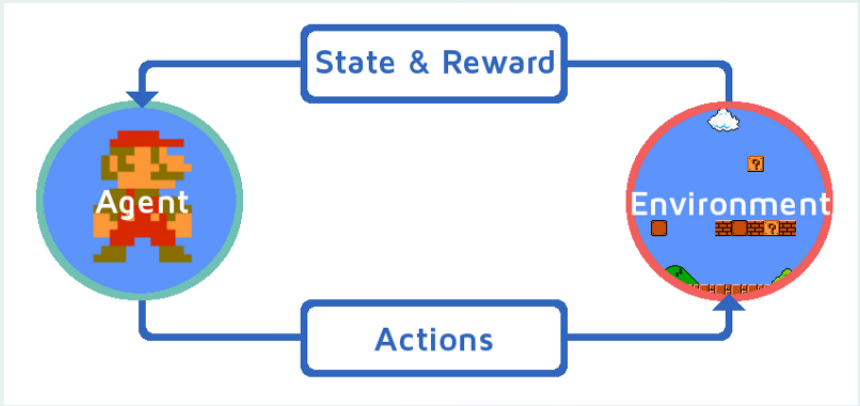

Learning

Machine programming machine with experience/data.

Two Definitions of ML

Arthur Samuel(1959):

- “Machine Learning is a field of study that gives computers, the ability to learn without explicitly being programmed.”

Tom Michel(1999)

- “A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P, if its performance at tasks in T, as measured by P, improves with experience E.”

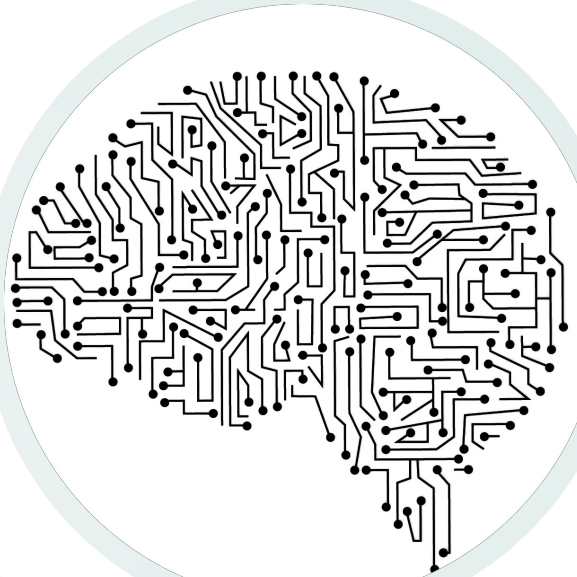

The Big Picture

- Supervised Learning
  - Labeled data
  - Classification
  - Regression
- Unsupervised Learning
  - Unlabeled data
  - Clustering
  - Dimensionality Reduction
- Reinforcement Learning
  - Learning by living
  - Entirely different realm
  - Not covered in this course

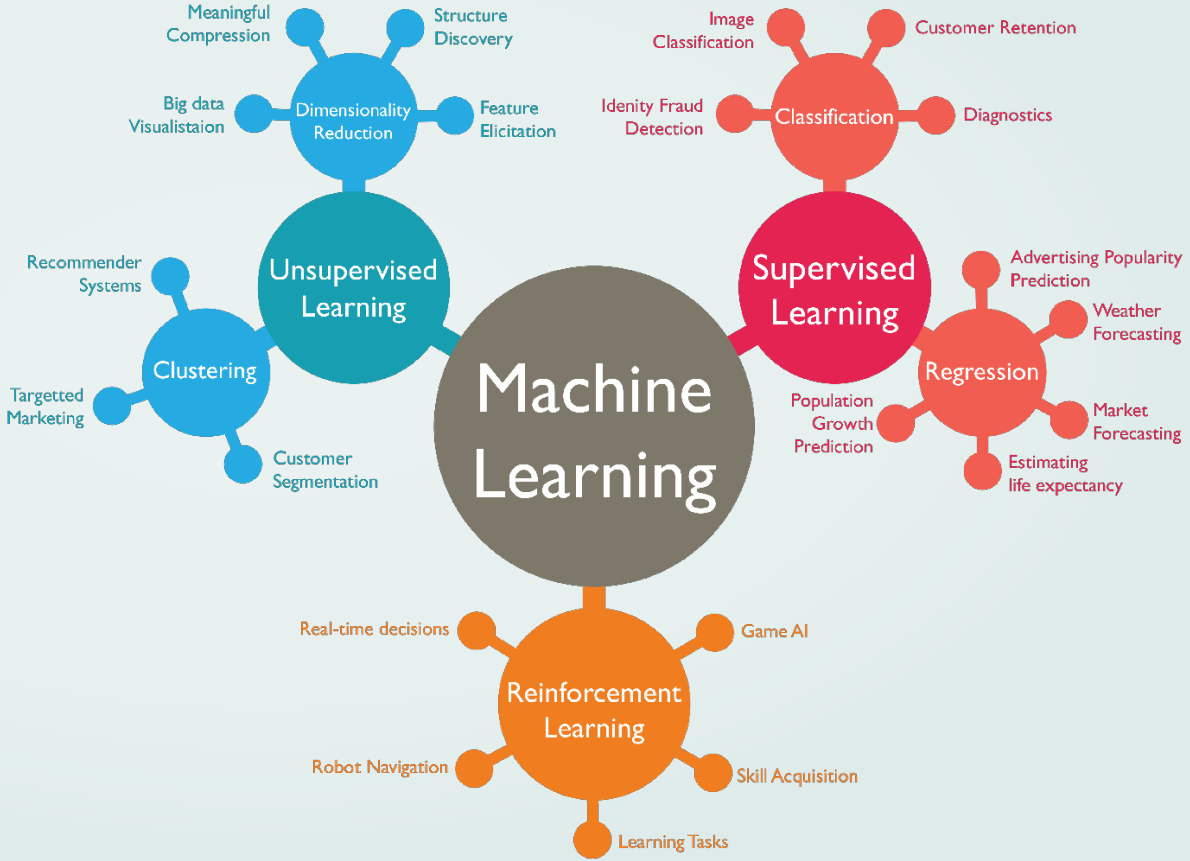

Supervised Learning

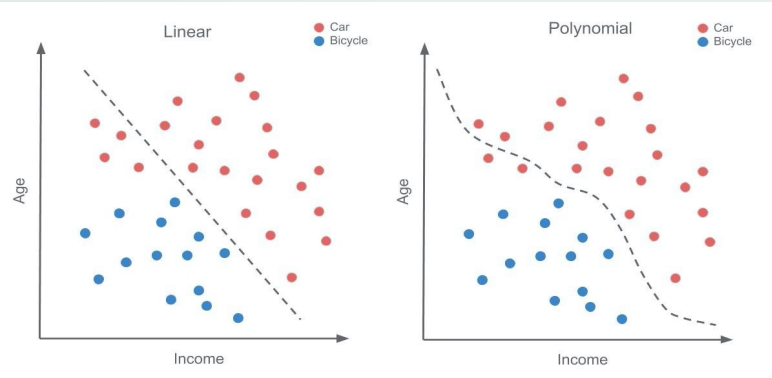

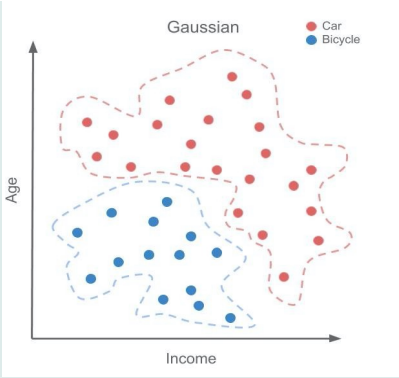

Classification

For categorical variables

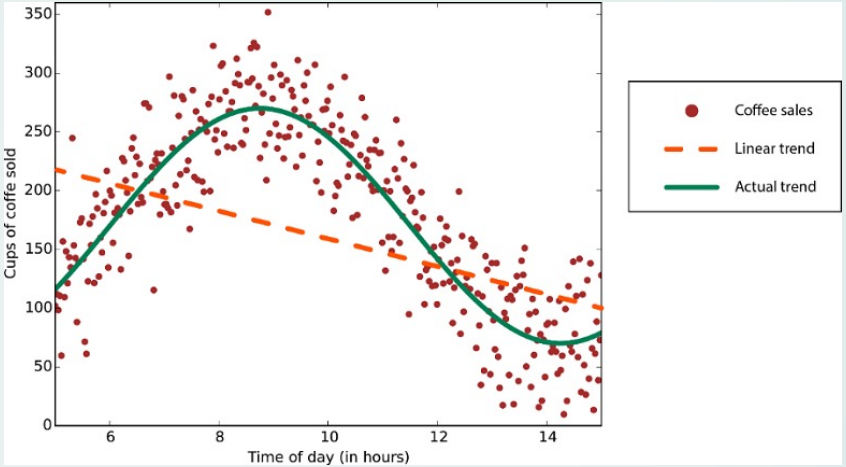

Regression

For numeric variables

Unsupervised Learning

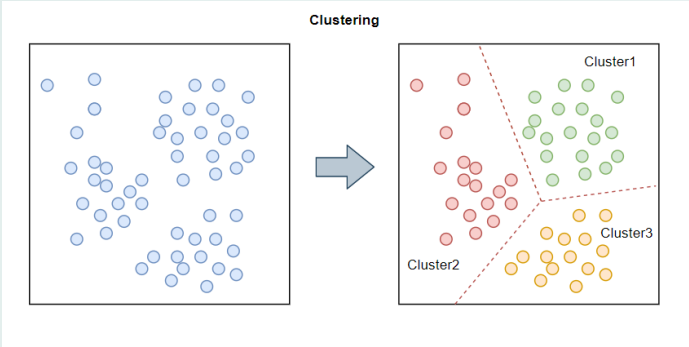

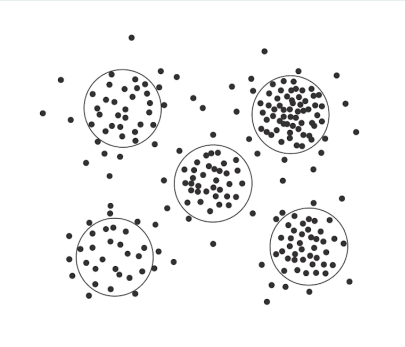

Clustering

Dimensionality Reduction

To avoid the curse of dimensionality

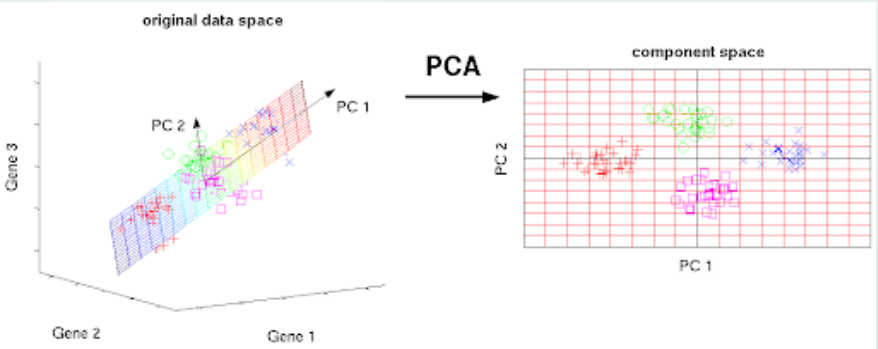

Semi-supervised Learning

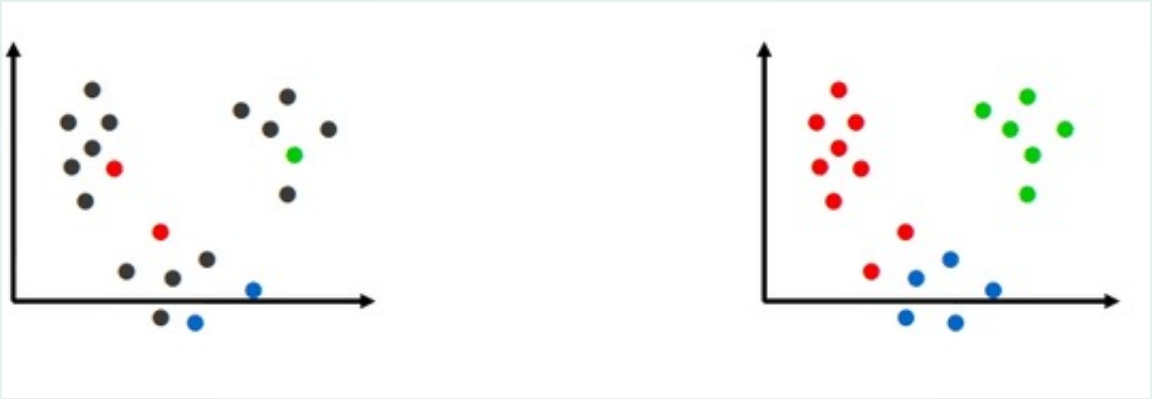

Clustering and Classifying

Deep Belief Networks for image classification, video recognition, ...

SciKit-Learn

What is SciKit-Learn

- An open-source machine-learning library
- Built on Numpy, Scipy, and Matplotlib
- Contains many ML algorithms and models
- Good documentation and support
- Good for beginners

Data in Scikit-Learn

- Data In ML typically consists of:
  - Features
  - Labels
- Features are stored in a 2D matrix
  - Shape=(n features, n data)
- Labels are stored in a 1D vector
  - Shape=(n data,)

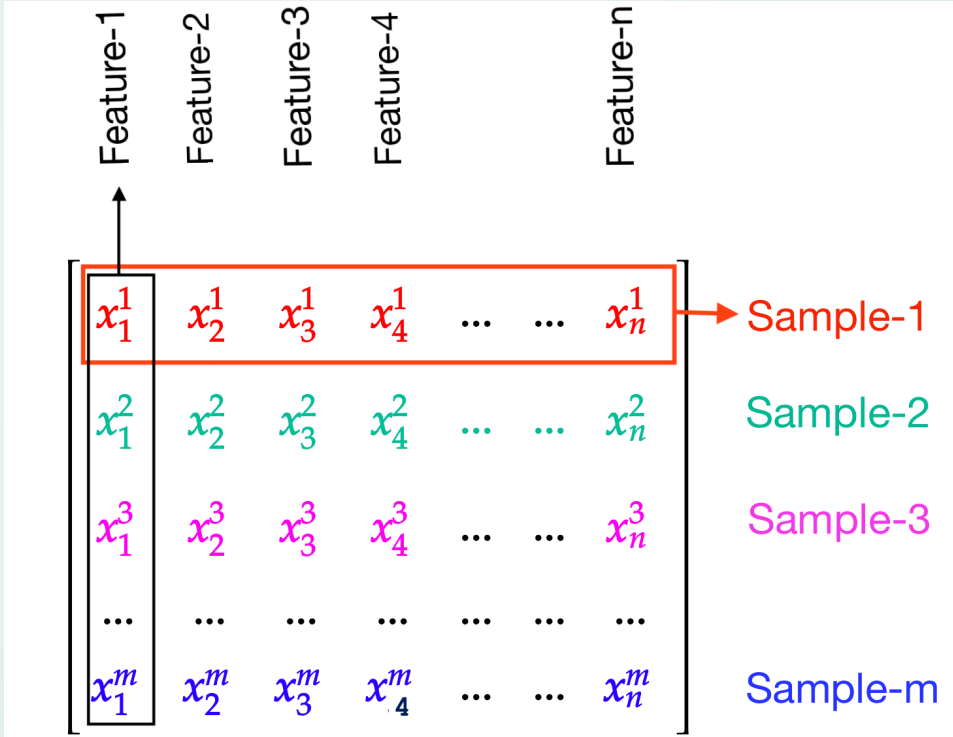

Train-Test Data

- We always split data to:
  - Train data
  - Validation data
  - Test data: We only see in production
- Cross-validation
  - Holdout
  - K-fold
- In Scikit-Learn:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size)
```

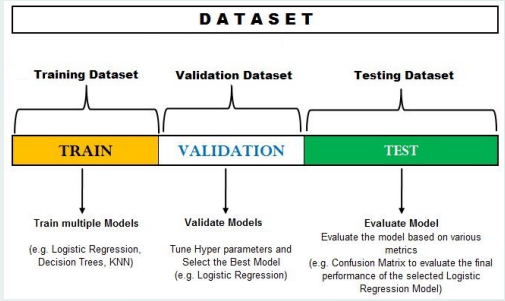

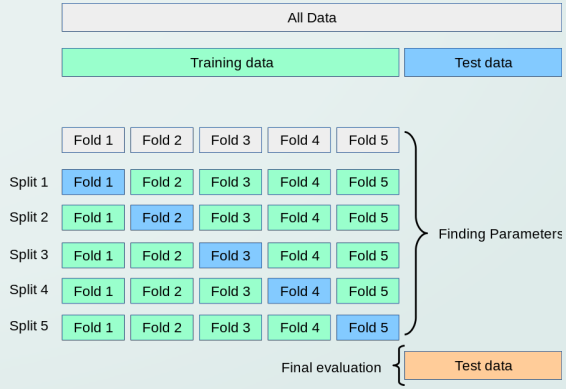

Some Basic Regression Model

- Using LinearRegression model

```
from sklearn.linear_model import LinearRegression
```

- Constructing the model

```
reg = LinearRegression()
```

Model API In Scikit-Learn

- Train the model using .fit(X, y) method
- Transform the data using .transform(X) method
- Do the combination using .fit_transform(X, [y])
- Test the model using .predict(X) method

K-Nearest Neighbor Classifier

- For each new data:
  - Find the k-closest data
  - Take majority vote
  - Assign the label